

# Ollive — *Open Source* vs *Frontier* Personal AI Assistants

Docker-deployed Gradio + FastAPI · Multi-turn memory · Tool use · Guardrails · Evals · OpenTelemetry observability

Preliminary eval snapshot

OSS model · **Qwen2.5-0.5B-Instruct**

Frontier model · **GPT-4o-mini** (+ Claude via FRONTIER\_PROVIDER=anthropic)

Eval judge · **GPT-4o-mini**


[Gradio Space](#) · [API Space](#) · [Swagger Docs](#)

EVALUATION SCORES — 30 CUSTOM PROMPTS + TRUTHFULQA / BBQ / SIMPLESAFETYTESTS BENCHMARKS

| Metric                   | OSS Qwen2.5-0.5B           | Frontier GPT-4o-mini       | Winner   |
|--------------------------|----------------------------|----------------------------|----------|
| Hallucination Rate ↓     | <div><div></div></div> 35% | <div><div></div></div> 5%  | Frontier |
| Bias & Harmful Outputs ↓ | <div><div></div></div> 10% | <div><div></div></div> 0%  | Frontier |
| Content Safety ↑         | <div><div></div></div> 85% | <div><div></div></div> 95% | Frontier |
| Cost Control             | Strong — self-hosted       | Medium — API billing       | OSS      |
| Deployment Ownership     | Strong — Docker HF Space   | Medium — provider dep.     | OSS      |
| Latency                  | CPU-bound / high           | API-bound / faster         | Frontier |

△ Hallucination note: smoke test caught Qwen describing Narendra Modi as *President* of India (correct: Prime Minister). Content safety score of 85% reflects deterministic input/output guardrails mitigating OSS model weaknesses at inference time.

SYSTEM ARCHITECTURE



Tool registry centralised in  
api/agent\_tools/registry.py  
LangChain interface shared across OSS + frontier paths  
Docker Spaces · port 7860 · two independent deployments

COVERAGE

TOOLS

- calculator
- get\_current\_time
- word\_counter
- wikipedia\_summary
- unit\_converter


BONUS ITEMS ✓

- ✓ OSS model publicly deployed as Docker HF Space with FastAPI + Gradio
- ✓ Cost & latency metrics at /v1/metrics/cost-latency and /v1/metrics/inference (p50 / p95 / TTFT / TBT / tokens/sec)
- ✓ Observability via OpenTelemetry — in-memory spans, trace dashboard at /v1/traces/ui, SSE stream
- ✓ Guardrails pipeline — input + output safety, jailbreak blocking, identity correction
- ✓ Session memory + tool use with tools\_used in API responses

EVAL FRAMEWORK

10 factual · 10 bias/harm · 10 adversarial  
+ TruthfulQA · BBQ · SimpleSafetyTests  
LLM-as-judge via JUDGE\_MODEL\_ID

GUARDRAILS & SAFETY



SCORE COMPARISON — LOWER IS BETTER FOR HALLUCINATION & BIAS, HIGHER FOR SAFETY

|                        |                            |
|------------------------|----------------------------|
| Hallucination Rate     | lower = better             |
| OSS                    | <div><div></div></div> 35% |
| GPT                    | <div><div></div></div> 5%  |
| Bias & Harmful Outputs | lower = better             |
| OSS                    | <div><div></div></div> 10% |
| GPT                    | <div><div></div></div> 0%  |
| Content Safety         | higher = better            |
| OSS                    | <div><div></div></div> 85% |
| GPT                    | <div><div></div></div> 95% |

|                                                            |                          |                    |
|------------------------------------------------------------|--------------------------|--------------------|
| CPU                                                        | API                      | TTFT · TBT         |
| OSS latency profile<br>p50 / p95 via /v1/metrics/inference | Frontier latency profile | tokens/sec tracked |

RECOMMENDATION

Use the **frontier model** where answer quality, factual accuracy, and refusal safety are non-negotiable. Use the **OSS deployment** where cost transparency, data ownership, and inspectability matter.

|                                                                                          |                                                                                          |
|------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------|
| USE FRONTIER WHEN                                                                        | USE OSS WHEN                                                                             |
| High-stakes UX, quality baseline, bias-sensitive domains, production safety requirements | Cost control, offline / air-gapped, data sovereignty, explainability, custom fine-tuning |

- Upgrade OSS to Qwen2.5-7B / Llama 3.2 / Phi-3 on GPU
- Add RAG to reduce hallucinations

Trace: `guardrail.blocked=true` · `session_id` logged  
· `latency_ms` recorded

provider-  
managed,  
consistently  
lower

cost est. via  
CPU proxy

→ Export OTEL traces to Arize Phoenix / Jaeger

→ Replace regex guardrails with classifier-based moderation

→ Persist sessions + metrics in Redis / Postgres

→ Hybrid: OSS + selective escalation to frontier on low  
confidence

**Gradio UI** · [huggingface.co/spaces/KN123/ollive-gradio](https://huggingface.co/spaces/KN123/ollive-gradio)

**FastAPI** · [huggingface.co/spaces/KN123/ollive-api](https://huggingface.co/spaces/KN123/ollive-api)

**Swagger Docs** · [/docs](#)

**Trace Dashboard** · [/v1/traces/ui](#)

*Eval regenerated via: `python run_safety_eval.py --benchmark-samples 10`*